



UNIVERSITÀ DEGLI STUDI DI FIRENZE

FACOLTÀ DI INGEGNERIA

Probabilistic topic models: Latent Dirichlet Allocation

Marco Righini

21 Febbraio 2013

Probabilistic topic models

Collective knowledge (scientific articles, books, images. . .) continues to be digitized and stored.

Growing interest in tools for understanding, organizing and searching information in these vast domains.

Topic models

Statistical models that analyze the words of documents to discover the **hidden thematic structure** and annotate the documents according to that, through an unsupervised approach.

Probabilistic topic models

Collective knowledge (scientific articles, books, images. . .) continues to be digitized and stored.

Growing interest in tools for understanding, organizing and searching information in these vast domains.

Topic models

Statistical models that analyze the words of documents to discover the **hidden thematic structure** and annotate the documents according to that, through an unsupervised approach.

Precursors of LDA

- TF-IDF [Salton et Gills, 1983] : identifies sets of words discriminative for a set of documents but doesn't reveal too much of inter or intra documents statistical relations;
- LSI [Deerwester et al., 1990] : based on SVD of TF-IDF matrix, derived features can capture some basic linguistic notions such as synonymy and polysemy;
- pLSI [Hofmann, 1999] : documents are reduced to a probability distribution on a fixed set of topics. Still there's no generative probabilistic model at level of documents.

Basic intuition

Assumptions of dimensionality reduction methods: exchangeability of words in a document and of documents in the corpus.

Theorem (De Finetti). *Any collection of infinitely exchangeable random variables has a representation as mixture distribution.*

Idea [Blei, Ng and Jordan, 2003]

Represent documents as random mixture over latent topics and each topic as a distribution over terms.

Basic intuition

Assumptions of dimensionality reduction methods: exchangeability of words in a document and of documents in the corpus.

Theorem (De Finetti). *Any collection of infinitely exchangeable random variables has a representation as mixture distribution.*


Idea [Blei, Ng and Jordan, 2003]

Represent documents as random mixture over latent topics and each topic as a distribution over terms.

Generative process

```

1: for  $k = 1 \rightarrow K$  do
2:   sample  $\phi_k \sim \text{Dirichlet}(\beta)$ 
3: end for
4: for  $d = 1 \rightarrow D$  do
5:   sample  $\theta_d \sim \text{Dirichlet}(\alpha)$ 
6:   for  $n = 1 \rightarrow N_d$  do
7:     sample  $z_{d,n} \sim \text{Multinomial}(\theta_d)$ 
8:     sample  $w_{d,n} \sim \text{Multinomial}(\phi_{z_{d,n}})$ 
9:   end for
10: end for
  
```



Dirichlet distribution

A Dirichlet distribution $\theta \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K)$ is a distribution over the $K-1$ simplex, i.e. vectors of K components that sum to one.

It has some important properties:

- **conjugate prior** of categorical and multinomial distributions;
- **mean** vector m with $m_i = \mathbb{E}[\theta_i | \alpha_1, \dots, \alpha_K] = \frac{\alpha_i}{\sum_{k=1}^K \alpha_k}$;
- **scale** $s = \sum_{k=1}^K \alpha_k$ controls the peakedness around the mean.

Typically **symmetric Dirichlet distribution** ($\alpha_i = \alpha$ for $i = 1 \rightarrow K$) is chosen when there is no prior knowledge favoring one component over another.

$\alpha = 1$ defines an uniform distribution, $\alpha \ll 1$ and $\alpha \gg 1$ respectively a sparser or more dense distribution.

Dirichlet distribution

A Dirichlet distribution $\theta \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K)$ is a distribution over the $K-1$ simplex, i.e. vectors of K components that sum to one.

It has some important properties:

- **conjugate prior** of categorical and multinomial distributions;
- **mean** vector m with $m_i = \mathbb{E}[\theta_i | \alpha_1, \dots, \alpha_K] = \frac{\alpha_i}{\sum_{k=1}^K \alpha_k}$;
- **scale** $s = \sum_{k=1}^K \alpha_k$ controls the peakedness around the mean.

Typically **symmetric Dirichlet distribution** ($\alpha_i = \alpha$ for $i = 1 \rightarrow K$) is chosen when there is no prior knowledge favoring one component over another.

$\alpha = 1$ defines an uniform distribution, $\alpha \ll 1$ and $\alpha \gg 1$ respectively a sparser or more dense distribution.

Approximate posterior inference

Joint distribution:

$$p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta) = \prod_{k=1}^K p(\phi_k | \beta) \prod_{d=1}^D p(\theta_d | \alpha) (\prod_{n=1}^{N_d} p(z_{d,n} | \theta_d) p(w_{d,n} | \phi_{z_{d,n}}))$$

Posterior distribution
$$p(\mathbf{z}, \theta, \phi | \mathbf{w}, \alpha, \beta) = \frac{p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta)}{p(\mathbf{w} | \alpha, \beta)}$$

Theoretically, the evidence can be obtained marginalizing the joint distribution over all possible configurations of the hidden variables. In practice, exact inference is intractable because the configurations domain is too large.

Some algorithms have been developed to efficiently **approximate** the posterior:

- *sampling-based algorithms* (Gibbs sampling);
- *variational algorithms*.

Approximate posterior inference

Joint distribution:

$$p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta) = \prod_{k=1}^K p(\phi_k | \beta) \prod_{d=1}^D p(\theta_d | \alpha) (\prod_{n=1}^{N_d} p(z_{d,n} | \theta_d) p(w_{d,n} | \phi_{z_{d,n}}))$$

Posterior distribution
$$p(\mathbf{z}, \theta, \phi | \mathbf{w}, \alpha, \beta) = \frac{p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta)}{p(\mathbf{w} | \alpha, \beta)}$$

Theoretically, the evidence can be obtained marginalizing the joint distribution over all possible configurations of the hidden variables. In practice, exact inference is intractable because the configurations domain is too large.

Some algorithms have been developed to efficiently **approximate** the posterior:

- *sampling-based algorithms* (Gibbs sampling);
- *variational algorithms*.

Approximate posterior inference

Joint distribution:

$$p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta) = \prod_{k=1}^K p(\phi_k | \beta) \prod_{d=1}^D p(\theta_d | \alpha) (\prod_{n=1}^{N_d} p(z_{d,n} | \theta_d) p(w_{d,n} | \phi_{z_{d,n}}))$$

Posterior distribution
$$p(\mathbf{z}, \theta, \phi | \mathbf{w}, \alpha, \beta) = \frac{p(\mathbf{w}, \mathbf{z}, \theta, \phi | \alpha, \beta)}{p(\mathbf{w} | \alpha, \beta)}$$

Theoretically, the evidence can be obtained marginalizing the joint distribution over all possible configurations of the hidden variables. In practice, exact inference is intractable because the configurations domain is too large.

Some algorithms have been developed to efficiently **approximate** the posterior:

- *sampling-based algorithms* (Gibbs sampling);
- *variational algorithms*.

Gibbs sampling

Gibbs sampling is an MCMC (Markov Chain Monte Carlo) algorithm based on the definition of a Markov Chain whose stationary distribution is the posterior of interest.

-
- 1: choose (e.g. at random) the initial values $x_1^{(0)}, x_2^{(0)}, \dots, x_p^{(0)}$ for the p latent variables
 - 2: **for** $t = 1 \rightarrow T$ (number of scans) **do**
 - 3: **for** $i = 1 \rightarrow p$ **do**
 - 4: sample $x_i^{(t)} \sim p(x_i | x_1^{(t)}, \dots, x_{i-1}^{(t)}, x_{i+1}^{(t-1)}, \dots, x_p^{(t-1)})$
 - 5: **end for**
 - 6: **end for**
-

Once the chain has “burned in”, the posterior can be approximated collecting samples at a certain lag (to reduce correlation) and averaging them.

Gibbs sampling

Gibbs sampling is an MCMC (Markov Chain Monte Carlo) algorithm based on the definition of a Markov Chain whose stationary distribution is the posterior of interest.

-
- 1: choose (e.g. at random) the initial values $x_1^{(0)}, x_2^{(0)}, \dots, x_p^{(0)}$ for the p latent variables
 - 2: **for** $t = 1 \rightarrow T$ (number of scans) **do**
 - 3: **for** $i = 1 \rightarrow p$ **do**
 - 4: sample $x_i^{(t)} \sim p(x_i | x_1^{(t)}, \dots, x_{i-1}^{(t)}, x_{i+1}^{(t-1)}, \dots, x_p^{(t-1)})$
 - 5: **end for**
 - 6: **end for**
-

Once the chain has “burned in”, the posterior can be approximated collecting samples at a certain lag (to reduce correlation) and averaging them.

Implementation overview


Development has embraced the following phases:

- project of the preprocessing module, including dataset format conversion and k-fold creation;
- project of the LDA module which cares about collapsed Gibbs sampling execution and perplexity computation;
- implementation.

Preprocessing: dataset format conversion

The preprocessing module carries out these tasks:

- **dataset format conversion**



Input dataset was obtained through tokenization, after removal of stop words and rare words.


Benefits from textual format

- not all terms in the vocabulary were present in the corpus;
- immediate view of documents content.

Preprocessing: dataset format conversion

The preprocessing module carries out these tasks:

- **dataset format conversion**



Input dataset was obtained through tokenization, after removal of stop words and rare words.

Benefits from textual format

- not all terms in the vocabulary were present in the corpus;
- immediate view of documents content.

Preprocessing: k-folds and usage

- **k-folds creation:** the documents are randomly shuffled to prevent possible sequential correlations in the original dataset. Afterwards, k train-validation pairs are generated.

Usage:

```
createKfold [-h] [--test_frac TEST_FRAC] [--k K] [--output OUTPUT] [--debug] corpus vocabulary
```

positional arguments:

```
corpus          path to the corpus file
vocabulary      path to the vocabulary file
```

optional arguments:

```
-h, --help          show this help message and exit
--test_frac TEST_FRAC fraction of corpus documents to retain as test set (default: 0)
--k K              number of folds (default: 1, no validation set)
--output OUTPUT    output files directory (default: same directory of input)
--debug           debug flag (default: false)
```

Preprocessing: k-folds and usage

- **k-folds creation:** the documents are randomly shuffled to prevent possible sequential correlations in the original dataset. Afterwards, k train-validation pairs are generated.

Usage:

```
createKFold [-h] [--test_frac TEST_FRAC] [--k K] [--output OUTPUT] [--debug] corpus vocabulary
```

positional arguments:

```
corpus          path to the corpus file
vocabulary      path to the vocabulary file
```

optional arguments:

```
-h, --help          show this help message and exit
--test_frac TEST_FRAC  fraction of corpus documents to retain as test set (default: 0)
--k K                number of folds (default: 1, no validation set)
--output OUTPUT      output files directory (default: same directory of input)
--debug              debug flag (default: false)
```

LDA module: collapsed Gibbs sampling

The LDA module performs two tasks:

- **collapsed Gibbs sampling**: obtained from standard Gibbs sampling marginalizing over θ and ϕ ;

Data: words $w \in$ documents d

Result: topic assignments z and counters $n_{d,k}$, $n_{k,w}$ and n_k

```

1 begin
2   randomly initialize  $z$  and increment counters;
3   foreach iteration do
4     foreach  $w \in w$  do
5        $topic \leftarrow z[w]$ ;
6        $n_{d,topic} += 1$ ;  $n_{topic,w} += 1$ ;  $n_{topic} += 1$ ;
7       for  $k = 0 \rightarrow K-1$  do
8          $p(z = k | \cdot) = (n_{d,k} + \alpha_k) \frac{n_{k,w} + \beta_w}{n_k + \beta \times V}$ ;
9       end
10       $topic \leftarrow$  sample from  $p(z | \cdot)$ ;
11       $z[w] \leftarrow topic$ ;
12       $n_{d,topic} += 1$ ;  $n_{topic,w} += 1$ ;  $n_{topic} += 1$ ;
13    end
14  end
15  return  $z$ ,  $n_{d,k}$ ,  $n_{k,w}$  and  $n_k$ 
16 end
```

$$\theta_{d,k} = \frac{n_{d,k} + \alpha_k}{\sum_{z=0}^{K-1} (n_{d,z} + \alpha_z)}$$

$$\phi_{k,w} = \frac{n_{k,w}}{\sum_{v=1}^{|V|} (n_{k,v} + \beta_v)}$$

LDA module: perplexity computation

- **perplexity computation:** perplexity is an indicator of the generalization performance of a model;

Given a new set of documents D_{new}

$$perpl(D_{new}) = \exp \left(- \frac{\log(p(\mathbf{w}_{new}|\phi))}{n_{new}} \right)$$

where

$$\begin{aligned} \log(p(\mathbf{w}_{new}|\phi)) &= \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log(p(w|\phi)) = \\ &= \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log \left(\sum_z p(w|z, \phi) p(z) \right) = \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log \left(\sum_{k=0}^{K-1} \phi_{k,w} \theta_{d,k} \right) \end{aligned}$$

θ for the new documents is obtained running collapsed Gibbs sampling on a new Markov Chain where the previously inferred counters $\mathbf{n}_{k,w}$ and \mathbf{n}_k (i.e. ϕ) are considered fixed.

LDA module: perplexity computation

- **perplexity computation:** perplexity is an indicator of the generalization performance of a model;

Given a new set of documents D_{new}

$$perpl(D_{new}) = \exp \left(- \frac{\log(p(\mathbf{w}_{new}|\phi))}{n_{new}} \right)$$

where

$$\begin{aligned} \log(p(\mathbf{w}_{new}|\phi)) &= \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log(p(w|\phi)) = \\ &= \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log \left(\sum_z p(w|z, \phi) p(z) \right) = \sum_{d=1}^{|D_{new}|} \sum_{w \in D_{new}^{(d)}} \log \left(\sum_{k=0}^{K-1} \phi_{k,w} \theta_{d,k} \right) \end{aligned}$$

θ for the new documents is obtained running collapsed Gibbs sampling on a new Markov Chain where the previously inferred counters $\mathbf{n}_{k,w}$ and \mathbf{n}_k (i.e. ϕ) are considered fixed.

LDA module: usage

Usage:

```
lda [-h] [--validation VALIDATION] [--output OUTPUT] [--topics TOPICS] [--alpha ALPHA] [--beta BETA]
[--iters ITERS] [--burnin BURNIN] [--savestep SAVESTEP] [--topwords TOPWORDS] [--perplexity]
[--export_complete] --train TRAIN
```

Option arguments:

```
-h [ --help ]           print usage messages

-t [ --train ] arg      path to training set file

-v [ --validation ] arg path to validation set file (default: none). Needed if perplexity flag setted
-o [ --output ] arg     output files directory (default: same directory of input)
--topics arg            number of topics (default: 100)
--alpha arg             alpha hyperparameter value (default: 0.05*averageDocumentLength/#topics)
--beta arg              beta hyperparameter value (default: 200/#wordsInVocabulary)
--iters arg             number of Gibbs sampling iterations (default: 1000)
--burnin arg            number of Gibbs sampling iterations considered burnin (default: iters/2)
--savestep arg          step at which LDA is saved to harddisk (default: (iters-burnin)/10)
--topwords arg          number of top words to show for each topic (default: 20)
--perplexity            exports only perplexity (default: false, export variables). Needs validation
--export_complete       exports all variables (default: false, export only topwords)
```

Execution time

In the beginning, LDA module was implemented in Python language.

At execution time, it turned out to be too much time consuming also vectorializing the operations (over 4 minutes per iteration in a medium level configuration).

Therefore, implementation moved to C++ language.

Timing: C++ vs Python

C++ implementation is approximately eleven time faster than Python one.

Execution time

In the beginning, LDA module was implemented in Python language.

At execution time, it turned out to be too much time consuming also vectorializing the operations (over 4 minutes per iteration in a medium level configuration).

Therefore, implementation moved to C++ language.

Timing: C++ vs Python

C++ implementation is approximately eleven time faster than Python one.

Tests overview

Three datasets were selected for tests:

- **KOS**: set of 3430 blog entries from `dailykos.com`, an american political blog. Total number of words is 467714;
- **NIPS**: set of 1500 papers from NIPS conferences, for a sum of 1932365 words;
- **ENRON**: set of 39861 emails from Enron Corp, for an aggregate of 6412172 words.

For every dataset, tests concerned the following steps:

- ① **models evaluation**: models, corresponding to different values of topic number, were defined and then 4-fold cross-validated;
- ② **test on the whole dataset**: model with best performance was qualitatively selected and next applied to the whole dataset.

Tests overview

Three datasets were selected for tests:

- **KOS**: set of 3430 blog entries from `dailykos.com`, an american political blog. Total number of words is 467714;
- **NIPS**: set of 1500 papers from NIPS conferences, for a sum of 1932365 words;
- **ENRON**: set of 39861 emails from Enron Corp, for an aggregate of 6412172 words.

For every dataset, tests concerned the following steps:

- 1 **models evaluation**: models, corresponding to different values of topic number, were defined and then 4-fold cross-validated;
- 2 **test on the whole dataset**: model with best performance was qualitatively selected and next applied to the whole dataset.

Tests: cross-validation

During the k -th step of 4-fold cross-validation, collapsed Gibbs sampling is executed on the k -th training set of documents.

At regular intervals, perplexity is computed on the k -th validation set.

Convergence assessment

Convergence of collapsed Gibbs sampling is assessed through the perplexity on the validation set.

Tests: cross-validation

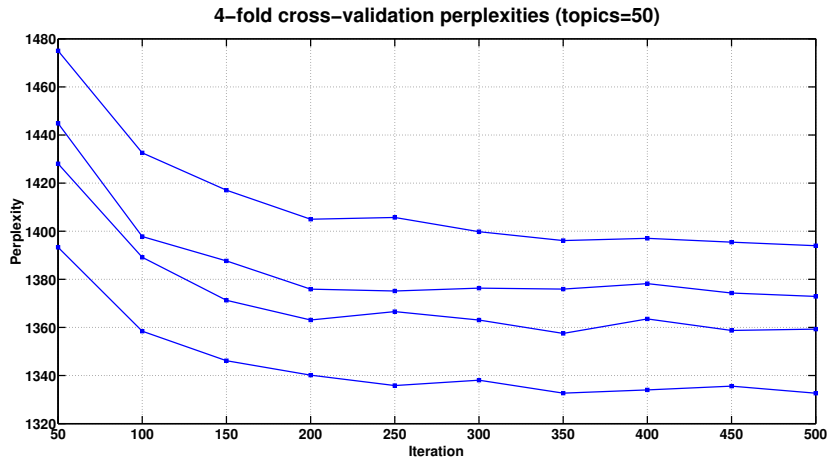
During the k -th step of 4-fold cross-validation, collapsed Gibbs sampling is executed on the k -th training set of documents.

At regular intervals, perplexity is computed on the k -th validation set.

Convergence assessment

Convergence of collapsed Gibbs sampling is assessed through the perplexity on the validation set.

KOS: perplexity



KOS: topics

million
money
democratic
campaign
party
dnc
raised
candidates
fundraising
national

kerry
john
edwards
democratic
general
presidential
kerrys
party
campaign
choice

november
voting
election
house
governor
account
electoral
republicans
senate
poll

military
abu
rumsfeld
ghraib
war
iraq
prisoners
prison
soldiers
abuse

iraq
war
weapons
bush
saddam
united
mass
threat
hussein
powell

cheney
bush
president
vice
dick
administration
general
lies
cheney's
truth

marriage
rights
gay
amendment
issue
law
abortion
ban
civil
constitution


investigation
law
federal
justice
documents
attorney
criminal
evidence
allegations
source

jobs
economy
economic
growth
job
rate
numbers
year
report
million

media
news
times
press
read
article
washington
papers
local
piece

NIPS: perplexity

4-fold cross-validation perplexities (topics=100)



NIPS: topics

network
neural
architecture
output
feedforward
feed
paper
architectures
forward
type

word
recognition
hmm
speech
system
training
context
mlp
continuous
hidden

kernel
vector
support
svm
function
set
margin
problem
examples
machines

face
human
subject
recognition
representation
faces
detection
facial
analysis
similarity

prior
bayesian
gaussian
posterior
distribution
mean
evidence
likelihood
covariance
mackay

processor
parallel
block
computer
memory
program
bit
performance
machine
operation

gradient
learning
weight
descent
convergence
rate
stochastic
algorithm
adaptive
perturbation


signal
frequency
channel
filter
power
low
processing
spectrum
detection
frequencies

cluster
clustering
set
rule
group
similarity
algorithm
partition
structure
number

coding
representation
code
bit
encoding
vector
codes
information
decoding
population

ENRON: perplexity

4-fold cross-validation perplexities (topics=50)



ENRON: topics

energy
program
plant
air
emission
environmental
demand
wind
generation
renewable

project
permit
station
facility
construction
site
water
unit
area
facilities

company
stock
financial
billion
dynegy
investor
shares
earning
analyst
trading

customer
cost
rate
contract
credit
rates
pay
amount
sce
period

page
court
labor
employees
law
worker
union
federal
employer
rules

access
user
password
center
account
web
link
help
message
site

student
business
school
program
university
haas
class
mba
event
berkeley

american
world
attack
country
government
war
article
bush
city
international

travel
hotel
roundtrip
fares
special
miles
city
ticket
offer
deal

free
click
online
offer
receive
link
special
web
visit
account

The project has led to:

- implementation of the LDA model and collapsed Gibbs Sampling;
- method, based on perplexity computation, for assessing the convergence of collapsed Gibbs Sampling;
- analysis of execution time between different programming languages (Python, C++);
- qualitative analysis of results obtained on three different corpus (KOS, NIPS and ENRON).

KOS: perplexities (mean of cross-validation) for different # topics

